

Step 1: Data Collection & Understanding



```
In [18]: # Importing pandas for data handling
import pandas as pd
# Pandas makes data cleaning and analysis faster and easier
# pandas can handle large datasets efficiently
class DataLoader:
    """
    This class is responsible for loading the dataset and providing an initial e

    Attributes:
        file_path (str): Path to the dataset file.
        df (DataFrame): Loaded dataset.
    """

    def __init__(self, file_path):
        """
        Initializes the DataLoader with the path to the dataset.

        Parameters:
            file_path (str): Path to the CSV dataset file.
        """
        self.file_path = file_path
        self.df = None

    def load_data(self):
        """
        Loads the dataset from the provided file path and prints important initi

        Returns:
            DataFrame: The loaded dataset.
        """
        print("\nDataset Loading...")
        print("-" * 100)

        # Load the dataset
        self.df = pd.read_csv(self.file_path)

        # Print shape of the dataset (rows, columns)
        print("\nDataset Loaded Successfully!")
        print("→ Shape of Dataset:", self.df.shape)
        print("-" * 100)

        # Print the data types of each column
        print(" Data Types:\n")
        print(self.df.dtypes)
        print("-" * 100)

        # Display first 5 rows of the dataset
        print("First 5 Rows:\n")
        print(self.df.head())
        print("-" * 100)

        # Display missing values in the dataset
```

```
print("Missing Values (Column-wise):\n")
print(self.df.isnull().sum())
print("-" * 100)

# Show statistical summary of numeric features
print("Statistical Summary (Numerical Columns):\n")
print(self.df.describe().T)
print("-" * 100)

return self.df

loader = DataLoader("heart.csv")
df = loader.load_data()
```

Dataset Loading...

Dataset Loaded Successfully!

→ Shape of Dataset: (918, 12)

Data Types:

```
Age          int64
Sex          object
ChestPainType  object
RestingBP    int64
Cholesterol  int64
FastingBS    int64
RestingECG   object
MaxHR        int64
ExerciseAngina object
Oldpeak      float64
ST_Slope     object
HeartDisease int64
dtype: object
```

First 5 Rows:

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	40	M	ATA	140	289	0	Normal	172	
1	49	F	NAP	160	180	0	Normal	156	
2	37	M	ATA	130	283	0	ST	98	
3	48	F	ASY	138	214	0	Normal	108	
4	54	M	NAP	150	195	0	Normal	122	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0

Missing Values (Column-wise):

```
Age          0
Sex          0
ChestPainType 0
RestingBP    0
Cholesterol  0
FastingBS    0
RestingECG   0
MaxHR        0
ExerciseAngina 0
Oldpeak      0
ST_Slope     0
HeartDisease 0
dtype: int64
```

Statistical Summary (Numerical Columns):

	count	mean	std	min	25%	50%	75%	max
Age	918.0	53.510893	9.432617	28.0	47.00	54.0	60.0	77.0
RestingBP	918.0	132.396514	18.514154	0.0	120.00	130.0	140.0	200.0
Cholesterol	918.0	198.799564	109.384145	0.0	173.25	223.0	267.0	603.0
FastingBS	918.0	0.233115	0.423046	0.0	0.00	0.0	0.0	1.0
MaxHR	918.0	136.809368	25.460334	60.0	120.00	138.0	156.0	202.0
Oldpeak	918.0	0.887364	1.066570	-2.6	0.00	0.6	1.5	6.2
HeartDisease	918.0	0.553377	0.497414	0.0	0.00	1.0	1.0	1.0

Step 2: Preprocessing, Univariate & Bivariate Analysis

```
In [19]: import seaborn as sns # Built on top of matplotlib, seaborn makes it easy to cre
import matplotlib.pyplot as plt # Low-level plotting library for creating visual
import numpy as np # Provides numerical operations and supports fast array
from sklearn.preprocessing import StandardScaler # From Scikit-learn, this class
# Crucial for machine learning algorithms that rely on feature scaling.

class DataPreprocessor:
    """
    This class handles preprocessing, univariate analysis, and bivariate analysis
    """

    def __init__(self, df):
        self.df = df

    def preprocess(self):
        """
        Handles missing values, encodes categorical variables, and returns the p
        """
        print("Starting Data Preprocessing...")
        print("-" * 100)

        # Handling missing values (if any)
        if self.df.isnull().sum().sum() > 0:
            self.df = self.df.fillna(self.df.median(numeric_only=True))
            print("Missing values filled with median of respective columns.")

        # Encoding categorical features
        self.df = pd.get_dummies(self.df, drop_first=True)
        print("Categorical Variables Encoded.")
        print("-" * 100)

        print("Data Preprocessing Completed Successfully.")
        print("-" * 100)
        return self.df

    def univariate_analysis(self):
        """
        Performs univariate analysis including countplots, pie charts, and histo
        """
        print("Univariate Analysis Started...")
        print("-" * 100)
```

```

# Countplot of Target Variable
sns.countplot(x='HeartDisease', data=self.df)
plt.title("Target Class Distribution (Countplot)")
plt.show()

# Pie Chart of Target Variable
target_counts = self.df['HeartDisease'].value_counts()
plt.figure(figsize=(6, 6))
plt.pie(target_counts, labels=['No Heart Disease', 'Heart Disease'], autopct='%1.1f%%')
plt.title("Target Class Distribution (Pie Chart)")
plt.show()

# Histograms of Numerical Features
numeric_cols = self.df.select_dtypes(include=[np.number]).columns.drop('HeartDisease')
self.df[numeric_cols].hist(bins=20, figsize=(14, 10), edgecolor='black')
plt.suptitle("Histogram of Numerical Features", fontsize=16)
plt.show()

print(self.df['HeartDisease'].value_counts())
print("-" * 100)

def bivariate_analysis(self):
    """
    Performs bivariate analysis including correlation heatmap and boxplots of numerical features vs target
    """
    print("Bivariate Analysis Started...")
    print("-" * 100)

    # Correlation Heatmap
    plt.figure(figsize=(12, 8))
    sns.heatmap(self.df.corr(), annot=True, cmap='coolwarm_r')
    plt.title("Correlation Heatmap")
    plt.show()

    # Boxplots for Numerical Features vs Target
    numeric_cols = self.df.select_dtypes(include=[np.number]).columns.drop('HeartDisease')
    for col in numeric_cols:
        plt.figure(figsize=(8, 5))
        sns.boxplot(x='HeartDisease', y=col, data=self.df)
        plt.title(f"{col} vs HeartDisease (Boxplot)")
        plt.show()

    # Optional Pairplot for smaller datasets (commented out)
    # sns.pairplot(self.df, hue='HeartDisease', diag_kind='hist')
    # plt.show()

    print("-" * 100)

preprocessor = DataPreprocessor(df)
df = preprocessor.preprocess()
preprocessor.univariate_analysis()
preprocessor.bivariate_analysis()

```

Starting Data Preprocessing...


Categorical Variables Encoded.

Data Preprocessing Completed Successfully.

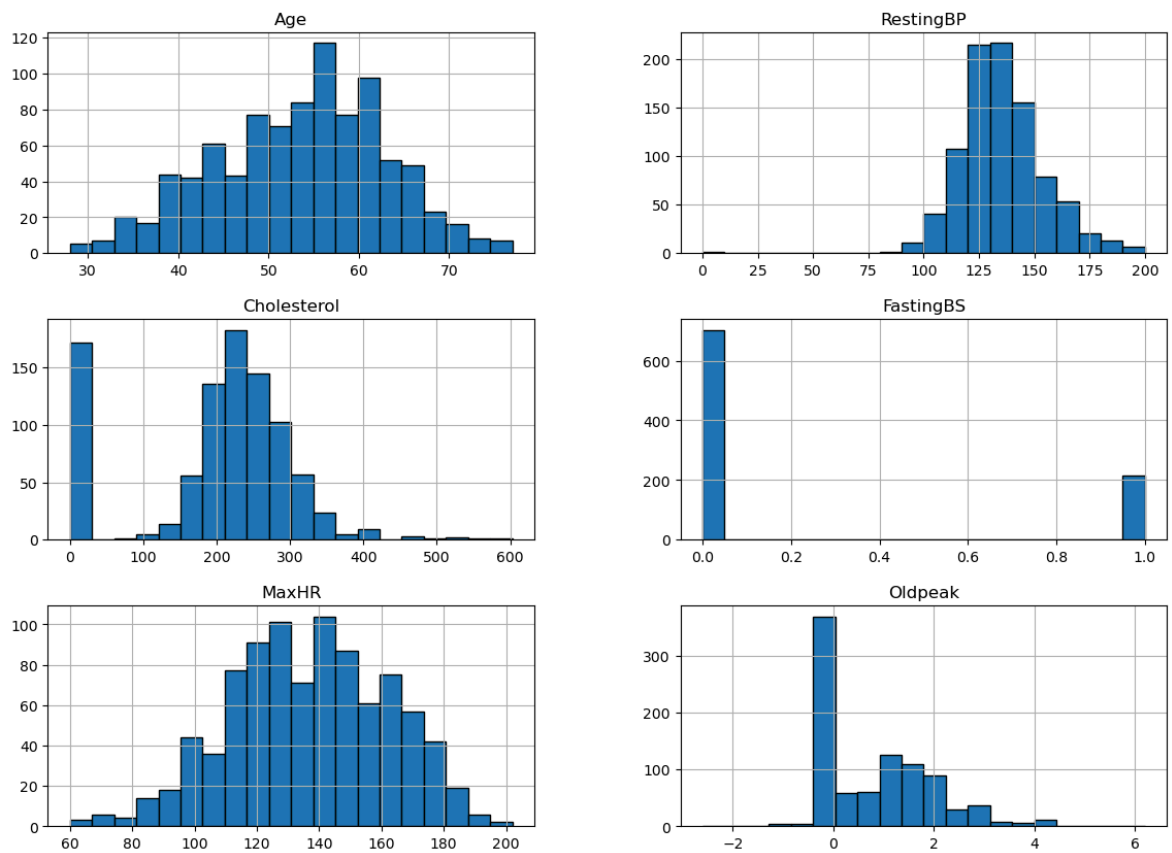
Univariate Analysis Started...



Target Class Distribution (Pie Chart)



Histogram of Numerical Features





HeartDisease



1 508


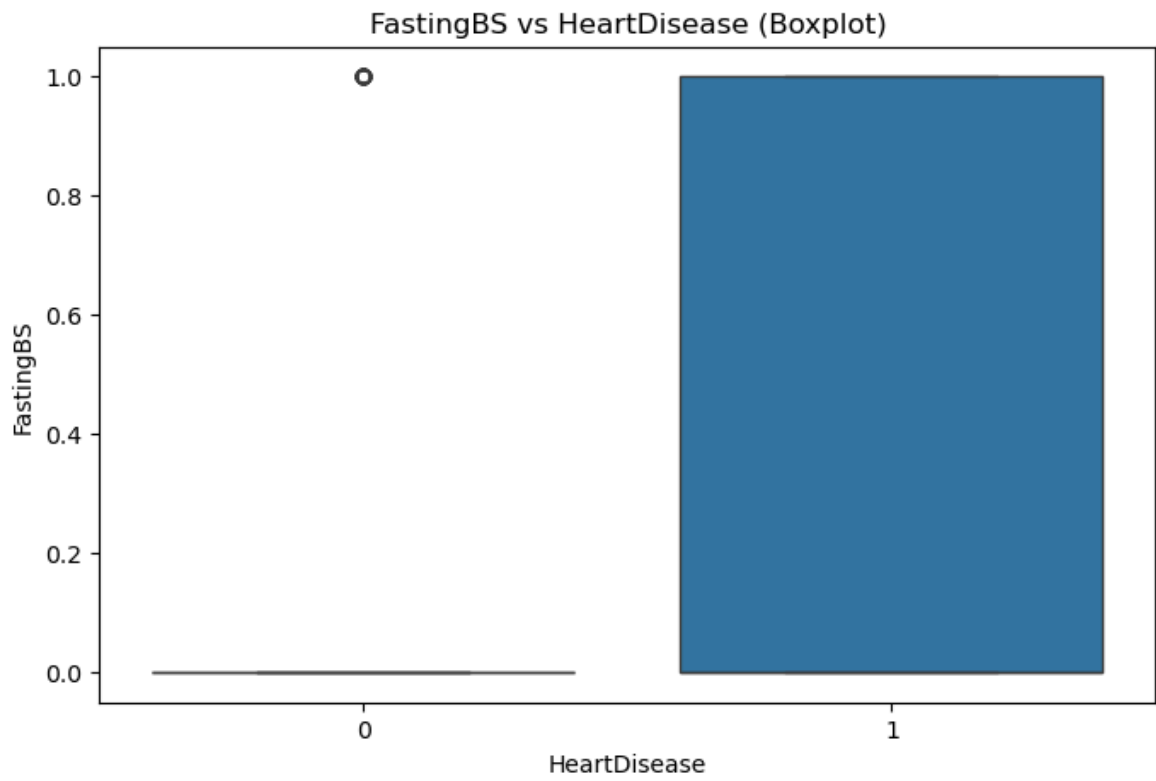
0 410


Name: count, dtype: int64

Bivariate Analysis Started...









Step 3: Data Splitting ✂

```
In [20]: from sklearn.model_selection import train_test_split # Splits your dataset into
from sklearn.preprocessing import StandardScaler # Standardizes/Scales your nume

class DataSplitter:
    """
    Splits the dataset into training and testing sets and applies feature scaling
    """

    def __init__(self, df):
        self.df = df

    def split(self):
        # Step 1: Separate Features and Target Variable
        X = self.df.drop('HeartDisease', axis=1)
        y = self.df['HeartDisease']

        # Step 2: Train-Test Split (70% training, 30% testing)
        X_train, X_test, y_train, y_test = train_test_split(
            X, y, test_size=0.3, random_state=42
        )

        # Step 3: Apply Feature Scaling
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)

        # Step 4: Display Shape of Data
        print("Data Split Completed:")
        print(f"Training Set Shape: {X_train.shape}")
```

```

print(f"Testing Set Shape: {X_test.shape}")
print("Data Scaled Successfully.")
print("-" * 100)

return X_train_scaled, X_test_scaled, y_train, y_test, scaler

# Example Usage
splitter = DataSplitter(df)
X_train, X_test, y_train, y_test, scaler = splitter.split()

```

Data Split Completed:

Training Set Shape: (642, 15)

Testing Set Shape: (276, 15)

Data Scaled Successfully.

Step 4: Model Training (SVM) and Evaluation

```

In [21]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report

class SVMTrainer:
    """
    Trains and evaluates a Support Vector Machine (SVM) classifier.
    """

    def __init__(self):
        self.model = None

    def train(self, X_train, y_train, kernel='linear'):
        """
        Train an SVM model on the provided training data.

        Parameters:
            X_train: Training features
            y_train: Training target labels
            kernel: Kernel type ('linear', 'rbf', 'poly', etc.)
        """
        self.model = SVC(kernel=kernel)
        self.model.fit(X_train, y_train)
        print("Model Training Completed.")
        print("-" * 100)
        return self.model

    def evaluate(self, model, X_test, y_test):
        """
        Evaluate the trained model on test data.

        Parameters:
            model: Trained model object
            X_test: Testing features
            y_test: Actual target labels for test data
        """
        print("Model Evaluation Results")
        print("-" * 100)
        y_pred = model.predict(X_test)

```

```

print(f"Accuracy Score: {accuracy_score(y_test, y_pred):.4f}")
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))
print("-" * 100)

# Example Usage
trainer = SVMTrainer()
model = trainer.train(X_train, y_train, kernel='linear')
trainer.evaluate(model, X_test, y_test)

```

Model Training Completed.

Model Evaluation Results

Accuracy Score: 0.8732

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.89	0.85	112
1	0.92	0.86	0.89	164
accuracy			0.87	276
macro avg	0.87	0.88	0.87	276
weighted avg	0.88	0.87	0.87	276

Step 5: Model Saving and Loading using Pickle

```

In [22]: import pickle

class ModelManager:
    """
    Handles saving and loading of trained ML models using pickle.
    """

    def save_model(self, model, scaler, filename='heart_disease_svm_model.pkl'):
        """
        Saves the trained model and scaler to a file.

        Parameters:
            model: Trained machine learning model
            scaler: Scaler object used for data preprocessing
            filename: Name of the file to save
        """
        with open(filename, 'wb') as f:
            pickle.dump((model, scaler), f)
        print(f"Model Saved as '{filename}'")
        print("-" * 100)

    def load_model(self, filename='heart_disease_svm_model.pkl'):
        """

```

Loads the trained model and scaler from a file.

Returns:

model: The trained machine learning model

scaler: The scaler object used for preprocessing

"""

with open(filename, 'rb') as f:

model, scaler = pickle.load(f)

print(f"Model Loaded from '{filename}')

return model, scaler

manager = ModelManager()

manager.save_model(model, scaler)

Saves model

loaded_model, loaded_scaler = manager.load_model()

Loads model back

print("Testing Reloaded Model Accuracy:")

print("-" * 100)

print("Score:", loaded_model.score(X_test, y_test))

Model Saved as 'heart_disease_svm_model.pkl'

Model Loaded from 'heart_disease_svm_model.pkl'

Testing Reloaded Model Accuracy:

Score: 0.8731884057971014